

Appendix R2: Foundations

Statistics and Regression Overview

J. Alberto Espinosa

2025-01-04

This Quarto document contains R code examples for basic statistical analysis associated with descriptive analytics. It also contains examples of predictive modeling with simple linear regression, bi-variate regression with a dummy variables and multivariate analysis. All predictive analytics work should start with very thorough descriptive analytics first, to become familiarized with the data at hand. For the corresponding conceptual material associated with these appendices, consult the main book [Predictive Analytics and Machine Learning for Managers](#). Chat with the PAML4M book and its related lecture slides and R scripts with its companion [PAML4M ChatBot](#).

Table of contents

Introduction	1
1. Covariance and Correlation (Quantitative x Quantitative)	2
2. Analysis of Variance (ANOVA) (Quantitative x Categorical)	11
The aov() Function	11
The anova() Function	17
3. Chi-Square Test of Independence (Categorical x Categorical)	18
4. Simple Linear Regression	20
Fitting a Model	20
The lm() Object	21
The summary() Object	25
Useful Regression Plots	26
Making Predictions With the Fitted Model	30
5. Bi-Variate Regression with a Dummy Variable	34
6. Multiple Linear Regression	35

This script was created by J. Alberto Espinosa for educational and training purposes. Feel free to use this material for your own work, but please do not share or duplicate without the author's permission.

Introduction

This Quarto document contains R code examples for basic statistical analysis associated with descriptive analytics. It also contains examples of predictive modeling with simple linear regression, bi-variate regression with a dummy variables and multivariate analysis. All predictive analytics work should start with very thorough descriptive analytics first, to become familiarized with the data at hand.

The most basic descriptive analytics often performed before start building predictive models include:

Visual inspection of the data, including things like scatter plots, distributions with histograms, QQ Plots, etc. Please refer to the PAML4M_1R_R_Intro.Qmd script.

Describing Central Tendencies and Variability

Descriptive Statistics, including things like means, medians, standard deviations, minimum/maximum values, outliers, etc.

Understanding Co-Variation

Covariance Analysis: (Quantitative x Quantitative) Covariance is an important and fundamental statistical concept to help understand if two variable co-vary in the same direction or opposite directions or not. But it is not very useful when variables have dissimilar scales because results will change if you change the scale of a variable (e.g., from inches to feet).

Correlation Analysis: (Quantitative x Quantitative) Is the covariance of two variables, divided by the standard deviation of each of the two variables. Because we divide by the variance of each variable, all issues of scale we have with covariance go away and the correlation values are bound between -1 (perfectly negatively correlated) to 0 (uncorrelated or independent) to +1 (perfectly positively correlated). Because correlation is based on differences in values, standard deviations and variance, you can only compute correlation when both variables involved are quantitative. If one of the variables (or both) is (are) binary you can still compute the correlation statistic, but it is not as useful as ANOVA.

ANOVA or Analysis of Variance: (Quantitative x Categorical) If one of the variables is quantitative and the other is categorical or even binary, it is more useful to evaluate co-variation with an ANOVA test. ANOVA computes the mean for the quantitative variable for each of the categories of the other (categorical) variable and evaluates if these means vary significantly across the categories. It is called analysis of “variance” and not analysis of “means” because the means are evaluated by comparing the variances in the two groups to see whether they vary more **within** each category (not significant or independent) than **across** categories (significant co-variation). I illustrate this more clearly below.

Chi-Square Test of Independence: (Categorical x Categorical) If you want to understand the co-variation between two categorical variables, correlation and ANOVA won’t help. To evaluate this, the typical approach is to cross tabulate all the categories of one variable in rows and all the categories of the other variables as columns, with the cross-tabulated counts in the respective cells. If the two variables are independent, the proportion of counts between any two cells in a given column (or row) will be similar to the proportion of counts for the respective row (or column) totals. But if one variable has a significant influence on the other, the cell proportions will be very different than the row (or column proportions). I illustrate this more clearly below

1. Covariance and Correlation (Quantitative x Quantitative)

It is important to develop a good sense for which variables co-vary or not with others. Generally speaking, when building predictive models, the goal is to have predictors that are highly correlated with the outcome (i.e., dependent) variable, but are not correlated with each other (i.e., independent variables). Thus, covariance and correlation matrices provide very useful information when developing predictive models.

Let’s look at the diamonds data in the ggplot package:

```
require(ggplot2) # Contains the diamonds data set
data(diamonds) # Load the data set into the work environment
```

The `cov()` and `cor()` functions in the `{stats}` package provide quick covariance and correlation, respectively. They both require a matrix as an input, so data frames need to be first converted into a matrix. For example, let’s bind 3 variable vectors into a data frame and convert it into a matrix:

```
MyDiamonds.dat <- data.frame(diamonds$price,
                             diamonds$carat,
                             diamonds$depth)
```

```
MyDiamonds.mat <- as.matrix(MyDiamonds.dat)
```

Alternatively, you can just specify the column numbers this way, all rows, columns 7, 1 and 5

```
MyDiamond.mat <- as.matrix(data.frame(diamonds[, c(7,1,5)]))
```

Before we continue, let's change the annoying scientific notation (the `scipen` = parameter stands for "scientific notation penalty"). For example, if we want only very small numbers in scientific location we can use the `options()` function with `scipen` = 4, so that only values with more than 4 zeros after the decimal point are displayed in scientific notation.

```
options(scipen="4")
```

Now you can compute the covariance and correlation matrices for these 3 variables

```
options(scipen = 4) # To limit the use of scientific notation
```

```
cov(MyDiamonds.mat, use = "complete.obs") # Discard rows w/incomplete data
```

	diamonds.price	diamonds.carat	diamonds.depth
diamonds.price	15915629.42430	1742.76536427	-60.85371214
diamonds.carat	1742.76536	0.22468666	0.01916653
diamonds.depth	-60.85371	0.01916653	2.05240384

```
cov(MyDiamonds.mat, use = "pairwise.complete.obs") # Discard pairs w/o data
```

	diamonds.price	diamonds.carat	diamonds.depth
diamonds.price	15915629.42430	1742.76536427	-60.85371214
diamonds.carat	1742.76536	0.22468666	0.01916653
diamonds.depth	-60.85371	0.01916653	2.05240384

```
cor(MyDiamonds.mat, use = "complete.obs") # Discard rows w/incomplete data
```

	diamonds.price	diamonds.carat	diamonds.depth
diamonds.price	1.0000000	0.92159130	-0.01064740
diamonds.carat	0.9215913	1.00000000	0.02822431
diamonds.depth	-0.0106474	0.02822431	1.00000000

```
cor(MyDiamonds.mat, use = "pairwise.complete.obs") # Discard pairs w/o data
```

	diamonds.price	diamonds.carat	diamonds.depth
diamonds.price	1.0000000	0.92159130	-0.01064740
diamonds.carat	0.9215913	1.00000000	0.02822431
diamonds.depth	-0.0106474	0.02822431	1.00000000

It is easy to figure out from the output above that the covariance matrix is difficult to interpret because it is affected by the scales of the variables involved. For example, the covariance of price with carats is 1742.76. What does this mean? Is this value high or low? Hard to tell. But if we look at the correlation matrix we see that the correlation is 0.92, which is close to 1, so yes, it is very high. Again, covariance is a very important statistical value of great

mathematical and modeling interest. There are many statistical methods that are based on the co-variance structure of the data. But for business interpretations, correlation is more useful. Consequently, we will use correlation from now on for the most part.

Unfortunately the `cor()` function only gives correlation values, not **p-values**. For p-values, the `rcorr()` function in the **{Hmisc}** package does the job. It returns the number of observations plus 2 matrices -> one with **correlation** values and one with the respective **p-values**.

```
library(Hmisc) # Load the library. Note the H is upper case

rcorr(MyDiamonds.mat, # Requires a matrix as an input
      type = "pearson")
```

	diamonds.price	diamonds.carat	diamonds.depth
diamonds.price	1.00	0.92	-0.01
diamonds.carat	0.92	1.00	0.03
diamonds.depth	-0.01	0.03	1.00

n= 53940


P

	diamonds.price	diamonds.carat	diamonds.depth
diamonds.price		0.0000	0.0134
diamonds.carat	0.0000		0.0000
diamonds.depth	0.0134	0.0000	

You can display correlations (above the diagonal), distributions (diagonal) and scatter plots (below the diagonal) all together with the `ggpairs()` function in the **{GGally}** library:

```
library(GGally) # Package with useful graphics displays

# WARNING -- the following takes a LONG TIME
ggpairs(MyDiamonds.dat) # It can be a bit slow
```





It works with categorical data too. Try this on your own. It takes a long time:

```
ggpairs(diamonds)
```

The `pairs()` function from the base `{graphics}` package works well too. It works with both, matrices and data frames:

```
pairs(MyDiamonds.dat)
```



It works with categorical data too. Try this on your own. It takes a long time:

```
pairs(diamonds)
```

Example with the **mtcars** data set:

```
library(Hmisc) # Load the library. Note the H is upper case
library(GGally) # Package with useful graphics displays
options(width = 100) # Wide width to avoid wrapping the correlation matrix
rcorr(as.matrix(mtcars))
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
mpg	1.00	-0.85	-0.85	-0.78	0.68	-0.87	0.42	0.66	0.60	0.48	-0.55
cyl	-0.85	1.00	0.90	0.83	-0.70	0.78	-0.59	-0.81	-0.52	-0.49	0.53
disp	-0.85	0.90	1.00	0.79	-0.71	0.89	-0.43	-0.71	-0.59	-0.56	0.39
hp	-0.78	0.83	0.79	1.00	-0.45	0.66	-0.71	-0.72	-0.24	-0.13	0.75
drat	0.68	-0.70	-0.71	-0.45	1.00	-0.71	0.09	0.44	0.71	0.70	-0.09
wt	-0.87	0.78	0.89	0.66	-0.71	1.00	-0.17	-0.55	-0.69	-0.58	0.43
qsec	0.42	-0.59	-0.43	-0.71	0.09	-0.17	1.00	0.74	-0.23	-0.21	-0.66
vs	0.66	-0.81	-0.71	-0.72	0.44	-0.55	0.74	1.00	0.17	0.21	-0.57
am	0.60	-0.52	-0.59	-0.24	0.71	-0.69	-0.23	0.17	1.00	0.79	0.06
gear	0.48	-0.49	-0.56	-0.13	0.70	-0.58	-0.21	0.21	0.79	1.00	0.27
carb	-0.55	0.53	0.39	0.75	-0.09	0.43	-0.66	-0.57	0.06	0.27	1.00

n= 32

P


	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
mpg											
cyl											
disp											
hp											
drat											
wt											
qsec											
vs											
am											
gear											
carb											

```


mpg      0.0000 0.0000 0.0000 0.0000 0.0000 0.0171 0.0000 0.0003 0.0054 0.0011
cyl 0.0000      0.0000 0.0000 0.0000 0.0000 0.0004 0.0000 0.0022 0.0042 0.0019
disp 0.0000 0.0000      0.0000 0.0000 0.0000 0.0131 0.0000 0.0004 0.0010 0.0253
hp 0.0000 0.0000 0.0000      0.0100 0.0000 0.0000 0.0000 0.1798 0.4930 0.0000
drat 0.0000 0.0000 0.0000 0.0100      0.0000 0.6196 0.0117 0.0000 0.0000 0.6212
wt 0.0000 0.0000 0.0000 0.0000 0.0000      0.3389 0.0010 0.0000 0.0005 0.0146
qsec 0.0171 0.0004 0.0131 0.0000 0.6196 0.3389      0.0000 0.2057 0.2425 0.0000
vs 0.0000 0.0000 0.0000 0.0000 0.0117 0.0010 0.0000      0.3570 0.2579 0.0007
am 0.0003 0.0022 0.0004 0.1798 0.0000 0.0000 0.2057 0.3570      0.0000 0.7545
gear 0.0054 0.0042 0.0010 0.4930 0.0000 0.0005 0.2425 0.2579 0.0000      0.1290
carb 0.0011 0.0019 0.0253 0.0000 0.6212 0.0146 0.0000 0.0007 0.7545 0.1290

```

```
ggpairs(mtcars) # Takes a bit of time
```




```
pairs(mtcars) # A bit quicker, but not as nice
```




The `corrplot()` function in the `{corrplot}` library provides nice graphics.

```
library(corrplot) # Library for correlation plots
mtCorr <- cor(mtcars) # First, store the correlation object


corrplot(mtCorr,
  method = "number",
  order = "hclust") # Show correlation
```



```
corrplot(mtCorr,
  method = "circle",
  order = "hclust") # Then plot it
```




```
corrplot(mtCorr,
  method = "ellipse",
  order = "hclust") # Slanted left/right for +/-
```



You can order variables clustered by correlation values and omit the diagonal

```
corrplot(mtCorr,
  method = "number",
  order = "hclust",
  diag = F,
  title = "MT Cars Correlation Matrix")
```



Type `?corrplot()` at the console to see all the methods available

2. Analysis of Variance (ANOVA) (Quantitative x Categorical)

Analysis of Variance (ANOVA) is a test that compares the means of 2 or more groups or categories. The means of 2 groups are always different, but are they different by very small amounts? Or are the differences substantial? This is what ANOVA tests. ANOVA tells us if the difference in means between or across categories is significant. Generally speaking, ANOVA compares the **within-category** variance for both groups against the **between-category** variance across both groups. If the between-group variance is significantly larger than the within-group variance, then the difference in means between the two groups is significant. Otherwise it is not.

ANOVA can be used any time you need to compare means between or across groups. For example, an OLS regression reports an R-Squared (i.e., explained variance) and a p-value for the entire regression model. The p-value is based on an ANOVA test, which analyzes whether the variance in the errors or residuals around the regression line are significantly smaller than the variance of the dependent variable values, relative to its overall mean. If the p-value or ANOVA test is significant, we say that the regression model has a significant explanatory power and provides significantly more explanation than just the mean and its variance (i.e., a model without predictors or the “null” model).

There are 2 popular functions to do ANOVA tests: `aov()` and `anova()`, both available in `{stats}`.

The `aov()` Function

The `aov()` function is useful to compare the means of a continuous variable (e.g., price) across various categories (e.g., clarity, color). In a nutshell, it tests whether the variance across groups is larger than within groups. The larger the F statistic, the more confidence we have that the variance across groups is significant. In the example below, the ANOVA test is significant at the $p \ll 0.05$ level.

```
require(ggplot2) # Contains the diamonds data set

aov(price ~ clarity,
     data = diamonds) # Run the ANOVA on a single factor
```

Call:

```
aov(formula = price ~ clarity, data = diamonds)
```

Terms:

	clarity	Residuals
Sum of Squares	23307802882	835165332636
Deg. of Freedom	7	53932

Residual standard error: 3935.165

Estimated effects may be unbalanced

The `summary()` of `aov()` provided more details about the F-test:

```
summary(aov(price ~ clarity,
             data = diamonds)) # More detailed results
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
clarity	7	23307802882	3329686126	215	<2e-16 ***
Residuals	53932	835165332636	15485525		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Alternatively, store the ANOVA results in an object

```
MyAOV <- aov(price ~ clarity,
             data = diamonds)

summary(MyAOV) # Show the ANOVA object result summary
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
clarity	7	23307802882	3329686126	215	<2e-16 ***
Residuals	53932	835165332636	15485525		

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

You can do ANOVA for a continuous variable across more than one categorical variable. In the examples below there are 2 and 3 ANOVA tests, respectively. The first shows that there is a significant difference in price across clarity categories, but analyzed within color. In the second test is the reverse, that is there is a difference in price across color, within clarity categories. The second ANOVA test works the same way but, for example, the difference in prices is significant across clarity categories, but within the same color and cut. And so on.

```
summary(aov(price ~ clarity + color,
             data = diamonds)) # ANOVA on 2 factors
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
clarity	7	23307802882	3329686126	222.4	<2e-16 ***
color	6	27663814802	4610635800	307.9	<2e-16 ***


```
Residuals    53926 807501517834    14974252
```

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
summary(aov(price ~ clarity + color + cut,
             data = diamonds)) # On 3 factors
```


	Df	Sum Sq	Mean Sq	F value	Pr(>F)	
clarity	7	23307802882	3329686126	223.89	<2e-16	***
color	6	27663814802	4610635800	310.02	<2e-16	***
cut	4	5574933424	1393733356	93.72	<2e-16	***
Residuals	53922	801926584410	14871974			

```
---
```


```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

You can also visualize the differences with boxplots. These visualizations should tell a similar story than the respective ANOVA tests, but without significance statistics, just visually:


```
boxplot(price ~ clarity,
         data = diamonds)
```



```
boxplot(price ~ cut,
         data = diamonds)
```



```
boxplot(price ~ clarity + cut,
        data = diamonds)
```



Replicating the ISLR textbook example with the credit Default data set

This is another illustration of ANOVA and it replicates an example provided in the ISLR book, but it also conducts an ANOVA test in addition to the plots. Similarly to the example in the ISLR book, I first extract a balanced sample from the data set. It is common in classification data sets to have unbalanced number of observations across categories. For example, medical data sets with disease diagnosis tend to have far more negatives than positives. The same is true for transaction fraud, etc. This example is no different. The **Default** data set is synthetic and contains 10,000 observations. Only 333 of them represent loan defaults. To balance the analysis I use all 333 default records, but extract a random sample of 333 records from the no-default group.

```
library(ISLR) # Contains the Default data set

set.seed(1)

def.yes <- subset(Default,
                  default == "Yes") # Subset with all default loans

def.no <- subset(Default,
                 default == "No") # Subset with all no-default loans

def.no <- def.no[sample(nrow(def.no), 333),] # Extract 333 of them randomly

def.sub <- rbind(def.yes, def.no) # Bind the defaults and no defaults


# Scatterplot with vertical lines at the respective means

plot(income ~ balance,
     data = def.sub,
     col = ifelse(default == "No", "darkgreen", "purple"),
     pch = ifelse(default == "No", 1, 3))


abline(v = mean(def.sub$balance),
       col = "red")

# abline(h=mean(def.sub$income), col="red")


abline(h = 40000, col = "red")
```



```
boxplot(balance ~ default,
        data = def.sub,
        col = c("darkgreen", "purple"))
```



```
boxplot(income ~ default,
        data = def.sub,
        col = c("darkgreen", "purple"))
```



```
summary(aov(balance ~ default,
             data = def.sub))
```

```
              Df    Sum Sq  Mean Sq F value Pr(>F)
default        1 143029929 143029929   850.7 <2e-16 ***
Residuals     664 111636896   168128
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
summary(aov(income ~ default,
             data = def.sub))
```

```
              Df    Sum Sq  Mean Sq F value Pr(>F)
default        1  277297260 277297260    1.535  0.216
Residuals     664 119976574966 180687613
```

The anova() Function

The `anova()` function is useful when comparing nested linear models. One model is nested within another if the larger model contains all the predictors of the small model. Let's fit 3 nested regression models with the diamonds data set and price as the outcome variable. The first model is the null model with no predictors (`~ 1` in the model

indicates the intercept as the only predictor, and in this case nothing else). The second model is a simple linear regression model with a single predictors, and the last has 2 predictors.

```
library(ggplot2)

lm.null <- lm(price ~ 1,
              data = diamonds) # Null model

lm.small <- lm(price ~ carat,
              data = diamonds) # Small model

lm.large <- lm(price ~ carat + clarity,
              data = diamonds) # Large model

anova(lm.null, lm.small, lm.large) # Compare 3 nested models
```

Analysis of Variance Table

```
Model 1: price ~ 1
Model 2: price ~ carat
Model 3: price ~ carat + clarity
  Res.Df    RSS Df Sum of Sq    F    Pr(>F)
1  53939 858473135517
2  53938 129345695398  1 729127440120 435639.9 < 2.2e-16 ***
3  53931  90263944583  7  39081750815   3335.8 < 2.2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

The output shows 3 models and 2 ANOVA test results. The first ANOVA test compares the second model to the first in the `anova()` function. The second ANOVA test compares the thirds model to the second. You can test as many nested models as you wish in one pass. The first p-value is significant, so the small model has significantly more explanatory power than the null model. The second p-value is also significant, so the large model has more explanatory power than the small model. In other words, carats explain more variance in price than just the mean price, but adding clarity as a predictor improves the explanatory power over the small model.

3. Chi-Square Test of Independence (Categorical x Categorical)

For example, say that we want to see if diamond “cut” and “color” co-vary. In other words, we want to know if diamond cut and color are independent (i.e. one does not affect the value of the other) or if they are dependent (i.e., the value of one variable influences the value of another).

The first step is to prepare a cross table with the counts for both categories

```
attach(diamonds)
cross.table <- table (cut, color) # Store results in a cross table
cross.table # Check it out
```

cut	color						
	D	E	F	G	H	I	J
Fair	163	224	312	314	303	175	119
Good	662	933	909	871	702	522	307
Very Good	1513	2400	2164	2299	1824	1204	678
Premium	1603	2337	2331	2924	2360	1428	808
Ideal	2834	3903	3826	4884	3115	2093	896

We can now look at the margin totals:

```
rowSums(cross.table) # Check the row totals
```

Fair	Good	Very Good	Premium	Ideal
1610	4906	12082	13791	21551

```
cat("\n") # Blank line
```

```
colSums(cross.table) # Check the column totals
```

D	E	F	G	H	I	J
6775	9797	9542	11292	8304	5422	2808

```
cat("\n") # Blank line
```

```
sum(rowSums(cross.table)) # Compute table totals
```

```
[1] 53940
```

```
sum(colSums(cross.table)) # Same result
```

```
[1] 53940
```

If the row and column variables are totally independent the proportion of `rowSums()` would be identical to the proportion of any two cells on the same two rows of the table. Similarly, the proportion `colSums()` would be identical to the proportion of any two cells on the same two columns.

In fact, you can create a table of expected values as follows: $\text{Exp.Cell}(i,j) = \text{rowSum}(i) * \text{colSum}(j) / \text{table.tot}$. If the **observed** (actual) values in the table are close to the **expected** values, the two variables are independent. If the observed values are very different than the expected values, the two variables are dependent (i.e., they co-vary)

The `chisq.test()` function does this test for you when you feed it the cross table. The null hypothesis is that the two variables are independent. A significant p-value rejects the null hypothesis suggesting that the variables are dependent. In addition, we can extract the observed and expected values from the `chisq.test()` object:

```
chiSq.diamonds <- chisq.test(cross.table) # Conduct the test
chiSq.diamonds$observed # Check the observed values, same as the cross.table
```

	color						
cut	D	E	F	G	H	I	J
Fair	163	224	312	314	303	175	119
Good	662	933	909	871	702	522	307
Very Good	1513	2400	2164	2299	1824	1204	678
Premium	1603	2337	2331	2924	2360	1428	808
Ideal	2834	3903	3826	4884	3115	2093	896

```
print(chiSq.diamonds$expected, digits = 2) # Check the expected values
```

	color						
cut	D	E	F	G	H	I	J
Fair	202	292	285	337	248	162	84
Good	616	891	868	1027	755	493	255
Very Good	1518	2194	2137	2529	1860	1214	629
Premium	1732	2505	2440	2887	2123	1386	718
Ideal	2707	3914	3812	4512	3318	2166	1122

```
chiSq.diamonds # Check out Chi Square test results
```

Pearson's Chi-squared test

```
data: cross.table
X-squared = 310.32, df = 24, p-value < 2.2e-16
```

The p-value is significant, so we reject the null hypothesis that cut and color are independent and conclude that they are dependent (i.e., they co-vary).

4. Simple Linear Regression

Let's work an example with the **Boston** housing dataset in the **{MASS}** package. We start by fitting a **null** model. Normally, you don't need to do this, but I do this here as a building block. Let's load the library. You may also want to type `?Boston` in the console to explore the variables in the dataset

```
library(MASS) # Contains the Boston data set
```

Fitting a Model

The `lm()` function in the **{stats}** library fits **linear models** based on the OLS regression model. The `lm()` model requires the **outcome variable**, followed by the `~` operator, followed by the **predictors**, separated from each other with a `+` sign, and the data attribute to indicate the dataset to be used for fitting the model. For the null model, we specify no predictors and use `1`, which is like having a column of 1's in the dataset, which yield the intercept. As before, **medv** is the median value of homes in the respective county in the Boston area in thousands of \$.

```
lm.fit.null <- lm(medv ~ 1,
                  data = Boston) # Use 1 for no predictors
summary(lm.fit.null) # The null model
```

Call:

```
lm(formula = medv ~ 1, data = Boston)
```

Residuals:

Min	1Q	Median	3Q	Max
-17.533	-5.508	-1.333	2.467	27.467

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	22.5328	0.4089	55.11	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 9.197 on 505 degrees of freedom


```
# The outcome variable mean and the null model intercept are the same
mean(Boston$medv)
```

```
[1] 22.53281
```

The null model is not very useful. We could obtain the same result just by calculating the mean of **medv**. Notice that there is no ANOVA F-test, model p-value or R-squared, because there is really nothing to report. In a simple regression model, we include a single predictor and expect to improve the explained variance of the model. The explained variance will most definitely increase when we add one predictor, but we need to figure out if the increase is significant or not.

A simple regression model has only one predictor so we don't use the $+$ symbol yet. For example, let's add **lstat** as a predictor. This variable contains the percentage of population with low income status in the county.

```
lm.fit <- lm(medv ~ lstat,
             data = Boston) # Using the fixed data set Boston
summary(lm.fit)
```

Call:

```
lm(formula = medv ~ lstat, data = Boston)
```

Residuals:

Min	1Q	Median	3Q	Max
-15.168	-3.990	-1.318	2.034	24.500

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	34.55384	0.56263	61.41	<2e-16 ***
lstat	-0.95005	0.03873	-24.53	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 6.216 on 504 degrees of freedom

Multiple R-squared: 0.5441, Adjusted R-squared: 0.5432

F-statistic: 601.6 on 1 and 504 DF, p-value: < 2.2e-16

Notice that we specify the data set in the formula. This is convenient because it allows us to run the same model with a different data set very quickly. But, alternatively, you could attach the data to the environment and you would no longer need to specify the data set in the formula:

```
attach(Boston)
lm.fit <- lm(medv ~ lstat)
```

The lm() Object

The `lm()` function will yield an **lm** object, which we are naming **lm.fit**. This object contains properties (i.e., data) and methods (i.e., functions) to help us extract information from the model and do things like plotting, testing, etc. Let's extract some information from **lm.fit**. Let's start by displaying the object itself.

```
lm.fit # Bare bones information
```

```
Call:
lm(formula = medv ~ lstat, data = Boston)
```

```
Coefficients:
(Intercept)      lstat
      34.55      -0.95
```

The `summary()` function always provide a summary of the most important properties stored in the model, which you would be expected to need in an analysis.

```
summary(lm.fit) # Display the summary of results
```

```
Call:
lm(formula = medv ~ lstat, data = Boston)
```

```
Residuals:
    Min       1Q   Median       3Q      Max
-15.168  -3.990  -1.318   2.034  24.500
```

```
Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)  34.55384    0.56263   61.41  <2e-16 ***
lstat        -0.95005    0.03873  -24.53  <2e-16 ***
---


```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 6.216 on 504 degrees of freedom
Multiple R-squared:  0.5441,    Adjusted R-squared:  0.5432
F-statistic: 601.6 on 1 and 504 DF,  p-value: < 2.2e-16
```

You can inspect the coefficients visually too with the `coefplot()` library and function. The dot represents the coefficient and the bands represent the confidence interval. If the confidence interval does not cross the 0 vertical line, then the coefficient is significant.

```
require(coefplot)
coefplot(lm.fit)
```



In the output above we can see all coefficients, their respective standard errors and p-values. We can also see all key fit statistics for the regression model. If we want more specific data, we can use various functions.

For example, the `names()` function gives you a list of properties contained in the object:

```
names(lm.fit) # List all the properties of the lm.fit object
```

```
[1] "coefficients" "residuals"    "effects"      "rank"         "fitted.values" "assign"
[7] "qr"           "df.residual"  "xlevels"      "call"         "terms"        "model"
```

Each of these contains useful data, for example the regression coefficients:

```
lm.fit$coefficients
```

```
(Intercept)    lstat
34.5538409   -0.9500494
```

The regression formula:

```
lm.fit$call # Regression formula
```

```
lm(formula = medv ~ lstat, data = Boston)
```

The first 12 residual values:

```
lm.fit$residuals[1:12] # First 12 residual values
```

1	2	3	4	5	6	7	8	9
-5.8225951	-4.2703898	3.9748580	1.6393042	6.7099222	-0.9040837	0.1552726	10.7396042	10.3811363
10	11	12						
0.5920031	-0.1253316	-3.0466860						

The first 12 predicted values:

```
lm.fit$fitted.values[1:12] # First 12 predicted values
```

1	2	3	4	5	6	7	8	9	10
29.822595	25.870390	30.725142	31.760696	29.490078	29.604084	22.744727	16.360396	6.118864	18.307997
11	12								
15.125332	21.946686								

You can also use functions to extract data, for example, the coefficients:

```
coef(lm.fit)
```

```
(Intercept)      lstat
34.5538409    -0.9500494
```

```
confint(lm.fit) # Confidence intervals for the coefficients
```

```

          2.5 %      97.5 %
(Intercept) 33.448457 35.6592247
lstat      -1.026148 -0.8739505
```

To inspect the object in more detail you can use the `str()` function to display the whole structure of the object:

```
str(lm.fit) # So, for example, use:
```

```

List of 12
 $ coefficients : Named num [1:2] 34.55 -0.95
  .. attr(*, "names")= chr [1:2] "(Intercept)" "lstat"
 $ residuals    : Named num [1:506] -5.82 -4.27 3.97 1.64 6.71 ...
  .. attr(*, "names")= chr [1:506] "1" "2" "3" "4" ...
 $ effects      : Named num [1:506] -506.86 -152.46 4.43 2.12 7.13 ...
  .. attr(*, "names")= chr [1:506] "(Intercept)" "lstat" "" "" ...
 $ rank         : int 2
 $ fitted.values: Named num [1:506] 29.8 25.9 30.7 31.8 29.5 ...
  .. attr(*, "names")= chr [1:506] "1" "2" "3" "4" ...
 $ assign       : int [1:2] 0 1
 $ qr           :List of 5
  ..$ qr        : num [1:506, 1:2] -22.4944 0.0445 0.0445 0.0445 0.0445 ...
  .. ..- attr(*, "dimnames")=List of 2
  .. .. ..$ : chr [1:506] "1" "2" "3" "4" ...
  .. .. ..$ : chr [1:2] "(Intercept)" "lstat"
  .. ..- attr(*, "assign")= int [1:2] 0 1
  ..$ qraux: num [1:2] 1.04 1.02
  ..$ pivot: int [1:2] 1 2
  ..$ tol   : num 0.0000001
  ..$ rank  : int 2
```

```

..- attr(*, "class")= chr "qr"
$ df.residual : int 504
$ xlevels      : Named list()
$ call         : language lm(formula = medv ~ lstat, data = Boston)
$ terms        :Classes 'terms', 'formula' language medv ~ lstat
.. ..- attr(*, "variables")= language list(medv, lstat)
.. ..- attr(*, "factors")= int [1:2, 1] 0 1
.. .. ..- attr(*, "dimnames")=List of 2
.. .. .. ..$ : chr [1:2] "medv" "lstat"
.. .. .. ..$ : chr "lstat"
.. ..- attr(*, "term.labels")= chr "lstat"
.. ..- attr(*, "order")= int 1
.. ..- attr(*, "intercept")= int 1
.. ..- attr(*, "response")= int 1
.. ..- attr(*, ".Environment")=<environment: R_GlobalEnv>
.. ..- attr(*, "predvars")= language list(medv, lstat)
.. ..- attr(*, "dataClasses")= Named chr [1:2] "numeric" "numeric"
.. .. ..- attr(*, "names")= chr [1:2] "medv" "lstat"
$ model         : 'data.frame': 506 obs. of 2 variables:
..$ medv : num [1:506] 24 21.6 34.7 33.4 36.2 28.7 22.9 27.1 16.5 18.9 ...
..$ lstat: num [1:506] 4.98 9.14 4.03 2.94 5.33 ...
..- attr(*, "terms")=Classes 'terms', 'formula' language medv ~ lstat
.. .. ..- attr(*, "variables")= language list(medv, lstat)
.. .. ..- attr(*, "factors")= int [1:2, 1] 0 1
.. .. .. ..- attr(*, "dimnames")=List of 2
.. .. .. .. ..$ : chr [1:2] "medv" "lstat"
.. .. .. .. ..$ : chr "lstat"
.. .. ..- attr(*, "term.labels")= chr "lstat"
.. .. ..- attr(*, "order")= int 1
.. .. ..- attr(*, "intercept")= int 1
.. .. ..- attr(*, "response")= int 1
.. .. ..- attr(*, ".Environment")=<environment: R_GlobalEnv>
.. .. ..- attr(*, "predvars")= language list(medv, lstat)
.. .. ..- attr(*, "dataClasses")= Named chr [1:2] "numeric" "numeric"
.. .. .. ..- attr(*, "names")= chr [1:2] "medv" "lstat"
- attr(*, "class")= chr "lm"

```

The summary() Object

Another interesting aspect is that when we create the `summary()` object, we can also extract data from this object, in addition to the data from `lm.fit` for example:

```

lm.sum <- summary(lm.fit) # Store the summary in an object
names(lm.sum) # Try str(lm.sum) on your own too

```

```

[1] "call"           "terms"          "residuals"      "coefficients"   "aliased"        "sigma"
[7] "df"             "r.squared"      "adj.r.squared"  "fstatistic"     "cov.unscaled"

```

Notice all the data we can extract from the `summary()` object, which we saved with the name `lm.sum`. For example:

```

lm.sum$r.squared

```

```

[1] 0.5441463

```

```
lm.sum$adj.r.squared # Adjusted R-Squared
```

```
[1] 0.5432418
```

```
lm.sum$coefficients["lstat","Pr(>|t|)"] # p-value for lstat
```

```
[1] 5.081103e-88
```

Notice that the `$coefficients` of **lm.sum** are different than those of **lm.fit**. The former is a data frame with coefficients, standard errors, p-values, etc. Let's say that you want to print standard errors for each coefficient. You just need to get the second column of the dataframe:

```
lm.sum$coefficients # The whole data frame
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	34.5538409	0.56262735	61.41515	3.743081e-236
lstat	-0.9500494	0.03873342	-24.52790	5.081103e-88

```
cat("\n")
```

```
lm.sum$coefficients[, 2] # Just the standard error in the second column
```

	lstat
(Intercept)	0.56262735
lstat	0.03873342

Now, let's do a little computational statistics and print a table with the actual 95% confidence interval, which is the coefficient ± 2 standard errors:

```
# Coefficient - 2 * Standard error
low.coef <- lm.sum$coefficients[, 1] - 2 * lm.sum$coefficients[, 2]

# Coefficient + 2 * Standard error
hi.coef <- lm.sum$coefficients[, 1] + 2 * lm.sum$coefficients[, 2]

cbind("From" = low.coef, "To" = hi.coef)
```


	From	To
(Intercept)	33.428586	35.6790956
lstat	-1.027516	-0.8725825

Useful Regression Plots

Residual Plots

The `plot()` function is what we call a **polymorphic** function. It sounds strange, but it means something simple. It means that the function will behave differently, depending what type of object we feed into it. For example, this function will simply yield a scatter plot:


```
plot(medv ~ lstat,  
     data = Boston)
```



To get general help on the `plot()` function, you can type `?plot` in the console. But since `plot()` is **polymorphic** you can get specific help for the type of object you want to plot. For example, to get help on `plot()` for `lm()` objects, type `?plot.lm`. In general, a `?function.` followed by the object type, will provide help specific to that object.

But the same function, when applied to an `lm()` object yields 4 key residual plots from the regression model. Notice that I subdivide the plot display into 4 panes with the `par()` function, otherwise I would get 4 large plots one after another. Please remember to reset the display to 1 row x 1 column:

```
par(mfrow = c(2, 2)) # Divide the output into a 2x2 frame by row  
plot(lm.fit) # Display 4 important regression diagnostic plots
```




```
par(mfrow = c(1, 1)) # Restore the output window
```

Inspecting regression plots is very important in OLS regression modeling because they provide useful information to evaluate the model assumptions:

1. Residual plot (inspect dispersion of residuals)
2. Residual qqplot (inspect normality of residuals)
3. Standardized residual (square root) plot – same as (1) but residuals are divided by their standard deviations
4. shows the influence of outliers – observations that fall outside the dotted lines (Cook's distance) have too much influence on the slope of the regression line.


If you are only interested in one of the plots, say the second one, you use `which = 2`:

```
plot(lm.fit, which = 2)
```


You can also inspect if the residuals are normally distributed with a histogram:

```
hist(lm.fit$residuals)
```



Plot the Data and Regression Line

```
plot(medv ~ lstat,  
     data = Boston) # Plot the (x,y) data  
  
abline(a = mean(Boston$medv),  
       b = 0,  
       lwd = 3) # Line at the mean of medv  
  
abline(lm.fit, col = "red") # Draw a red regression line through the plot
```



You can change the color of the line (e.g., `col = "red"`), the thickness (e.g., `lwd = 3`), the type (e.g., `lty = "dotted"`), dot size (e.g., `pch = 20`), or symbol to use instead of a dot (e.g., `pch = "+"`). Try these on your own. But note that you can only visualize regression lines with a single predictor.

Making Predictions With the Fitted Model

The `predict()` function displays predicted values from a model for all data points. I'm just displaying the first 12 below:

```
predict(lm.fit)[1:12]
```

```
1      2      3      4      5      6      7      8      9     10  
29.822595 25.870390 30.725142 31.760696 29.490078 29.604084 22.744727 16.360396  6.118864 18.307997  
11     12  
15.125332 21.946686
```

If you also want to display the lower and upper bounds of the respective 95% confidence intervals for these predictions, use (note that we use a comma at the end to get all 3 columns of the `predict()` object:

```
predict(lm.fit,
        interval = "confidence")[1:12, ]
```

	fit	lwr	upr
1	29.822595	29.025299	30.619891
2	25.870390	25.265246	26.475534
3	30.725142	29.873477	31.576807
4	31.760696	30.843594	32.677798
5	29.490078	28.712077	30.268079
6	29.604084	28.819515	30.388652
7	22.744727	22.201573	23.287882
8	16.360396	15.626114	17.094677
9	6.118864	4.696433	7.541295
10	18.307997	17.668272	18.947722
11	15.125332	14.321106	15.929557
12	21.946686	21.401770	22.491601

If you want predicted values with confidence intervals for 3 new data points for `lstat` of, say 5, 10 and 15, you have 2 options:

```
predict(lm.fit,
        data.frame(lstat = (c(5, 10, 15))),
        interval = "confidence")
```

	fit	lwr	upr
1	29.80359	29.00741	30.59978
2	25.05335	24.47413	25.63256
3	20.30310	19.73159	20.87461

```
cat("\n")
```

```
predict(lm.fit,
        data.frame(lstat = (c(5, 10, 15))),
        interval = "prediction")
```

	fit	lwr	upr
1	29.80359	17.565675	42.04151
2	25.05335	12.827626	37.27907
3	20.30310	8.077742	32.52846

`interval = "confidence"` provides the range of values representing a 95% probability (by default) that the model will predict values within that interval. If the interval does not include 0, then the coefficients are significant at the $p < 0.05$ level. This is your confidence interval of what the model is likely to predict.

To get other confidence levels, e.g. 99%, use the `level` = attribute:

```
predict(lm.fit,
        data.frame(lstat = (c(5, 10, 15))),
        interval = "confidence",
```

```
level = 0.99)
```

```
      fit      lwr      upr
1 29.80359 28.75578 30.85141
2 25.05335 24.29107 25.81562
3 20.30310 19.55096 21.05524
```

`interval = "prediction"` works a little different. It provides the 95% probability that an **actual prediction** will fall in the interval range. This is your confidence that an actual prediction will fall in that range. Generally, prediction intervals are larger than confidence intervals because they need to account for the variance in the outcome variable.

Technical note on fitting and predicting models. For now, we are using the full data set (e.g., Boston) to fit the model and to make predictions. Later on, we will split the data set into training (e.g., Boston[train,]) and test (e.g., Boston[-train,]) subsamples, and then fit or train the model with the training subsample and test it with the test subsample. This is called **cross-validation** which is central to **machine learning**.

Quick Simple Regression Examples

Boston Housing Data Set

```
library(MASS) # Contains the Boston dataset
lm.fit <- lm(medv ~ rm,
             data = Boston)
summary(lm.fit)
```

Call:

```
lm(formula = medv ~ rm, data = Boston)
```

Residuals:

```
      Min       1Q   Median       3Q      Max
-23.346  -2.547   0.090   2.986  39.433
```

Coefficients:

```
      Estimate Std. Error t value Pr(>|t|)
(Intercept) -34.671      2.650  -13.08  <2e-16 ***
rm           9.102      0.419   21.72  <2e-16 ***
---

```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 6.616 on 504 degrees of freedom


Multiple R-squared: 0.4835, Adjusted R-squared: 0.4825

F-statistic: 471.8 on 1 and 504 DF, p-value: < 2.2e-16

Let's do some plotting:

```
plot(medv ~ rm,
     data = Boston,
     main = "Regression for Boston House Values",
     xlab = "No Rooms", ylab = "Median Value $000")
abline(lm.fit, col = "red" )
```

Regression for Boston House Values



Hitters Baseball Player Salary Dataset

```
library(ISLR) # Contains the Hitter dataset
Hitters <- na.omit(Hitters) # Need to remove several omitted values
lm.fit <- lm(Salary ~ HmRun,
             data = Hitters) # Using home runs to predict salaries
summary(lm.fit) # Home runs has a positive significant effect
```

Call:

```
lm(formula = Salary ~ HmRun, data = Hitters)
```

Residuals:

Min	1Q	Median	3Q	Max
-748.73	-275.49	-79.27	184.72	1829.00

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	330.594	43.551	7.591	5.64e-13 ***
HmRun	17.671	2.995	5.900	1.13e-08 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 424.6 on 261 degrees of freedom


Multiple R-squared: 0.1177, Adjusted R-squared: 0.1143

F-statistic: 34.81 on 1 and 261 DF, p-value: 0.00000001125

The results show that home runs have a significant effect on players' salaries. For each additional home run, on average, a player's salary goes up by \$17,671.

Now, let's plot the regression, first the data points

```
plot(Salary ~ HmRun,
     data = Hitters,
     cex = 0.5,
     xlab = "X (Predictor)",
     ylab = "Y (Outcome)")
abline(a = mean(Hitters$Salary),
       b = 0) # Line at the mean of medv
abline(lm.fit, col = "red") # Draw the regression line on the plot
```



5. Bi-Variate Regression with a Dummy Variable

In the **Carseats** dataset in the **{ISLR}** library, **Price** is a continuous variable and **US** is a binary variable (**Yes** = 1, **No** = 0). You can inspect the dataset by typing `?Carseats` on the console.

```
library(ISLR) # Contains the "Carseats" data set

lm.fit <- lm(Sales ~ Price + US,
             data = Carseats)

summary(lm.fit)
```

Call:

```
lm(formula = Sales ~ Price + US, data = Carseats)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-6.9269	-1.6286	-0.0574	1.5766	7.0515

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	13.03079	0.63098	20.652	< 2e-16 ***
Price	-0.05448	0.00523	-10.416	< 2e-16 ***
USYes	1.19964	0.25846	4.641	0.00000471 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.469 on 397 degrees of freedom

Multiple R-squared: 0.2393, Adjusted R-squared: 0.2354

F-statistic: 62.43 on 2 and 397 DF, p-value: < 2.2e-16

The results suggest that, on average and holding car seat origin constant, for every \$1 increase in price, the total sales in the average location goes down by 54 units. Similarly, on average and holding price constant, a US car seat sells 1,199 more units than a non-US car seat.

Technical note: The binary variable **US** has two possible categorical (i.e., factor) values, **Yes** or **No**. When you model a categorical variable like this, R quantifies the model by creating two dummy variables by combining the variable name along with the respective categorical variables. In this case, it creates one variable called **USNo** with a value of 1 if the car seat is not from the US, and 0 otherwise. It also creates another variable called **USYes** with a value of 1 if the car seat is from the US, and 0 otherwise. Because **USNo** and **USYes** are perfectly linearly related (i.e., if one is 1 the other one has to be 0), it drops the first one alphabetically (i.e., **USNo**) from the model, which is why we see a predictor named **USYes**, not **US**. The intercept captures the effect of **USNo** (i.e., **USYes** = 0) when all predictors are 0. And the coefficient for **USYes** captures the difference in sales for US car seats, compared to non-US car seats.

6. Multiple Linear Regression

Just about everything discussed up to this applies to multivariate regressions. We just need to refine the wording in the interpretation a bit. Suppose that, as before, we want to predict median values of houses in Boston counties. But this time we anticipate that **lstat** may not be the only predictor, but that crime rate (**crim**), proportion of homes build prior to 1940 and proximity to the Charles River (**Chas**, 0 or 1) may have an affect on home values. The model syntax is the same as for simple linear models, but now we have more than one predictor, separated with the **+** operator

```
lm.fit <- lm(medv ~ lstat + crim + age + chas,
             data = Boston)
summary(lm.fit) # Display regression output summary
```

Call:

```
lm(formula = medv ~ lstat + crim + age + chas, data = Boston)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-15.594	-3.834	-1.319	1.932	24.224

Coefficients:


	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	32.73813	0.73635	44.460	< 2e-16 ***
lstat	-0.97132	0.05026	-19.326	< 2e-16 ***
crim	-0.07492	0.03543	-2.115	0.0350 *

```
age          0.02987    0.01220    2.448    0.0147 *
chas         4.44525    1.07516    4.135 0.0000417 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Residual standard error: 6.051 on 501 degrees of freedom
Multiple R-squared: 0.5706, Adjusted R-squared: 0.5672
F-statistic: 166.4 on 4 and 501 DF, p-value: < 2.2e-16

You can inspect the coefficients visually too:

```
library(coefplot)
coefplot(lm.fit)
```



To run a multiple regression on all available variables use the dot “.”. Also note that we are including the parameter `correlation = T` to display a correlation matrix along with the regression output.

```
lm.fit <- lm(medv ~ .,
             data = Boston)
summary(lm.fit, correlation = T)
```

Call:
lm(formula = medv ~ ., data = Boston)

Residuals:

Min	1Q	Median	3Q	Max
-15.595	-2.730	-0.518	1.777	26.199

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	36.4594884	5.1034588	7.144	3.28e-12	***
crim	-0.1080114	0.0328650	-3.287	0.001087	**
zn	0.0464205	0.0137275	3.382	0.000778	***
indus	0.0205586	0.0614957	0.334	0.738288	
chas	2.6867338	0.8615798	3.118	0.001925	**
nox	-17.7666112	3.8197437	-4.651	4.25e-06	***
rm	3.8098652	0.4179253	9.116	< 2e-16	***
age	0.0006922	0.0132098	0.052	0.958229	
dis	-1.4755668	0.1994547	-7.398	6.01e-13	***
rad	0.3060495	0.0663464	4.613	5.07e-06	***
tax	-0.0123346	0.0037605	-3.280	0.001112	**
ptratio	-0.9527472	0.1308268	-7.283	1.31e-12	***
black	0.0093117	0.0026860	3.467	0.000573	***
lstat	-0.5247584	0.0507153	-10.347	< 2e-16	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.745 on 492 degrees of freedom


Multiple R-squared: 0.7406, Adjusted R-squared: 0.7338

F-statistic: 108.1 on 13 and 492 DF, p-value: < 2.2e-16

Correlation of Coefficients:

	(Intercept)	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	black
crim	-0.06												
zn	-0.02	-0.09											
indus	0.06	0.04	0.11										
chas	-0.03	0.05	-0.01	-0.10									
nox	-0.55	0.06	0.04	-0.26	-0.03								
rm	-0.71	0.03	-0.16	0.09	-0.03	0.09							
age	0.08	0.00	0.12	0.00	-0.05	-0.27	-0.21						
dis	-0.36	0.12	-0.40	0.21	0.02	0.28	0.13	0.29					
rad	0.28	-0.27	0.11	0.28	-0.11	-0.14	-0.16	0.08	0.02				
tax	-0.11	0.01	-0.22	-0.44	0.12	-0.07	0.07	-0.03	-0.02	-0.79			
ptratio	-0.60	0.02	0.31	-0.13	0.09	0.33	0.16	-0.08	-0.09	-0.19	-0.08		
black	-0.29	0.12	-0.01	0.03	-0.05	0.07	0.10	-0.06	0.02	0.07	0.03	-0.03	
lstat	-0.30	-0.15	-0.05	-0.08	0.06	-0.07	0.53	-0.34	-0.04	-0.04	0.02	-0.05	0.16

```
coefplot(lm.fit)
```



The respective p-values show that predictors of interest like **chas**, **lstat**, etc. are significant. The effect of **lstat** can now be interpreted as follows: on average and **holding all other variables constant**, when the population of lower income status increases by one percentage point, the median value of houses are \$524 lower. The one thing different in this interpretation is the term **holding all other variables constant**, which basically says that the effect of **lstat** is independent of the other predictors because we are **controlling for** them.

Naturally, we can access the properties inside **lm.fit** and extract useful information, such as coefficients and confidence intervals:

```
coef(lm.fit) # Display only the regression coefficients
```

(Intercept)	crim	zn	indus	chas	nox
36.4594883851	-0.1080113578	0.0464204584	0.0205586264	2.6867338193	-17.7666112283
rm	age	dis	rad	tax	ptratio
3.8098652068	0.0006922246	-1.4755668456	0.3060494790	-0.0123345939	-0.9527472317
black	lstat				
0.0093116833	-0.5247583779				

```
cat("\n")
```





```
confint(lm.fit) # Display only the confidence intervals
```

	2.5 %	97.5 %
(Intercept)	26.432226009	46.486750761
crim	-0.172584412	-0.043438304
zn	0.019448778	0.073392139
indus	-0.100267941	0.141385193

chas	0.993904193	4.379563446
nox	-25.271633564	-10.261588893
rm	2.988726773	4.631003640
age	-0.025262320	0.026646769
dis	-1.867454981	-1.083678710
rad	0.175692169	0.436406789
tax	-0.019723286	-0.004945902
ptratio	-1.209795296	-0.695699168
black	0.004034306	0.014589060
lstat	-0.624403622	-0.425113133

We can also inspect the residual plots:

```
par(mfrow = c(2, 2))
plot(lm.fit)
```



```
par(mfrow = c(1, 1))
```