

Appendix R7B: Dimensionality - Dimension Reduction

Principal Components and Partial Least Squares Regressions

J. Alberto Espinosa

2024-06-22

This Quarto document contains R code examples about predictive modeling methods based on dimension reduction – Principal Components Regression (PCR) and Partial Least Squares Regression (PLSR). These models aim to correct issues of dimensionality, such as severe multicollinearity. In contrast to Ridge and LASSO, which shrink all coefficients to different lambda values, PCR and PLSR decompose the predictors into components, with each component being a linear weighted combination of all predictors. In general, the first component in either method is the one that captures the largest amount of variance in the predictors. The second component captures the second largest variance, and so on. If the first M components capture a large portion of the P predictor's variance (e.g., 70%) and $M \ll P$, then we would accomplish a substantial amount of dimension reduction. It is important to note that as we add more components, the model becomes closer to the OLS model. That is, a model with few components has more bias and less variance, thus less useful for interpretation, whereas a model with lots of components has less bias but more variance. Again, Cross-Validation can help us evaluate the optimal number of components to use. For the corresponding conceptual material associated with these appendices, consult the main book [Predictive Analytics and Machine Learning for Managers](#).

Table of contents

Dimensionality	2
Principal Components Regression (PCR)	3
PCR Loadings and Scores	5
PCR Coefficients	6
Predictions with PCR	8
Partial Least Squares (PLS) Regression	9
PLSR Loadings and Scores	11
PLSR Coefficients	12
Predictions with PLSR	14
Logistic and GLM Regressions with PCR and PLSR	15

This script was created by J. Alberto Espinosa for educational and training purposes. Feel free to use this material for your own work, but please do not share or duplicate without the author's permission.

Technical Note: All procedures involving Machine Learning and Cross-Validation are based on random sampling. Therefore, results may change when you run the analysis multiple times, depending on the kind of random number generator **RNGkind** and the seed selected. The results in this script may not match the outputs in the main text of the book. Follow the chapter and this script appendix independently. Let's first set the random number generator default

```
RNGkind(sample.kind = "default") # To use the R default RNG
```

Dimensionality

As discussed previously, large and complex models generally suffer from **dimensionality** issues, causing them to be over-fitting and have high variance. One of the most pervasive dimensionality issues is multicollinearity. **Regularized** methods like Ridge and LASSO help us reduce the detrimental effects of multicollinearity by shrinking (i.e. biasing) the predictors. In contrast, dimension reduction methods don't just help correct for multicollinearity, but actually **take advantage of the multicollinearity** structure of the data. In fact, dimension reduction methods like PCR and PLSR don't help much if the multicollinearity condition in the data is low. In contrast, But it can improve things dramatically when there is severe multicollinearity. This family of methods is particularly useful when there are many predictors relative to the number of observations.

The basic idea is this. If two predictors are uncorrelated, a scatter plot of the two variables will show a somewhat spherical cloud of data points with no clear alignment in one direction or another. In contrast, if two predictors are highly correlated, the scatter plot of these two predictors will show a thin cloud of data points aligned in one direction. The direction of this alignment is the direction in which the variance in the data is the highest. When there are more than 2 predictors, there will be more than one direction of alignment, depending on which variable is correlated with which. If you move the origin of the plot to the mean of all variables and then rotate the axes in the direction of highest variance, second highest, etc. you would be aligning the rotated axes with **1st. Principal Component (1st PC), 2nd PC**, etc.

The idea behind dimension reduction is that, for **N predictors**, there are **N PCs** and these PC's have some interesting properties:

1. They are perpendicular, i.e., **orthogonal** to each other. Once you have identify the 1st. PC, the 2nd PC can be found by taking an axis that is perpendicular to the 1st PC and rotating it around it until the direction of second largest variance is found. Same thing with the 3rd PC, and so on.
2. The resulting PCs are thus constructed as linear combinations of the predictors are independent or uncorrelated (i.e., a linear weighted sum of a portion of each predictor). Consequently, all predictors are represented to some extent in each PC. Also, because the PC's are rotated orthogonally (i.e., at right angles), the PC's are uncorrelated, so they are truly independent variables.
3. The PCs are ordered from highest to lowest variance. This allows you to explore the PCs to find the first **M PCs** that explain a substantial amount of variance in the data.
4. When the predictors are highly correlated, the first few PCs will explain a large proportion of variance in the data, and the last PCs will explain very little. In contrast, when there is little correlation among the predictors, you will need lots of components, and perhaps all of them to represent the variance of the data.
5. So, the goal is to identify first M PC's that explain, say 70% or 80% of the variance of the predictors and disregard the remaining PCs. If $M \ll P$ you will be achieving substantial dimension reduction, because you can now run a regression with M PC's, rather than with P predictors.
6. PCR and PLSR will estimate a 1 PC model, 2 PC model, etc., up to a P PC model that includes all PCs. The model (either PCR or PLSR) with all P PCs will yield identical results to the OLS or GLM model. As the number of PCs in the model decreases, the variance of the model decreases too, but the bias increases. Therefore, models with **more PCs** are **better for interpretations**, whereas models with the **lowest CV test deviance** are **best for predictions**.

There are a few dimension reduction methods, but two of the most popular ones are: (1) **Principal Components Regression (PCR)** in which the predictors are rotated to find the PCs, without taking into account whether these dimensions help predict the outcome variable (i.e., **unsupervised** method); and (2) **Partial Least Squares Regression (PLSR)**, which is similar to PCR, but the axes are further rotated to improve their correlation with the response variable (i.e., **supervised** method). PCR and PLSR are competing methods of similar kind and there is no guarantee that one method will be better than the other. It is recommended to test both and select the one with best results. Depending on the data, you never know which of the two models will yield better results.

Principal Components Regression (PCR)

I will use the `{pls}` R package for both PCR and PLSR and also the **Hitters** data set from the `{ISLR}` R package, containing baseball player salaries and 19 predictors. Note that the `pcr()` function syntax is similar to the `lm()` function syntax, with a few additional parameters. Let's start by loading the libraries and eliminate rows with omitted data.


```
library(pls) # Has the Principal Components Regression pcr() function
library(ISLR) # Has the Hitters data set
Hitters <- na.omit(Hitters) # Let's remove missing values
```

Predict Salary with all predictors. `scale = T` standardizes the predictors, which is recommended, especially when variables are in different scales (e.g., lbs, feet, etc.). Also, `validation = "CV"` does 10FCV and `validation = "LOO"` does LOOCV. We use all predictors in the model.

```
set.seed(2) # Arbitrary
pcr.fit <- pcr(Salary ~ .,
               data = Hitters,
               scale = T,
               validation = "CV")
```

The first step is to analyze the model results visually. We do this by plotting the **Scree Plot**, which shows the MSE or RMSE for all the PC models. I'm showing the plot with **MSE** but you can use change the `val.type =` parameter to `"RMSEP"` if you prefer the RMSE. Both graphs are similar, except for the scale of the error.

```
validationplot(pcr.fit,
               val.type = "MSEP",
               legendpos = "topright") # Plots the MSE
```



In the Scree Plot above, one can see that the MSE drops substantially from 0 PCs (Null model) to 1 PC and it then the graph **elbows** to the right. A **sharp elbow to the right** is a point of interest because it is a point where we achieve sharp error reductions, with little change afterwards. There is another sharp elbow to the right at 6 PCs and then another one at 16 PCs. These are all models of interest worth exploring further quantitatively. There is an elbow pointing upwards at 13 PCs, but this is not a point of interest, because the MSE turns downwards after that, so it is better to continue going to the right.

Let's analyze the model quantitatively.

```
summary(pcr.fit) # Take a look
```

```
Data:   X dimension: 263 19
      Y dimension: 263 1
Fit method: svdpc
Number of components considered: 19
```

VALIDATION: RMSEP

Cross-validated using 10 random segments.

	(Intercept)	1 comps	2 comps	3 comps	4 comps	5 comps	6 comps
CV	452	351.9	353.2	355.0	352.8	348.4	343.6
adjCV	452	351.6	352.7	354.4	352.1	347.6	342.7

	7 comps	8 comps	9 comps	10 comps	11 comps	12 comps	13 comps
CV	345.5	347.7	349.6	351.4	352.1	353.5	358.2
adjCV	344.7	346.7	348.5	350.1	350.7	352.0	356.5

	14 comps	15 comps	16 comps	17 comps	18 comps	19 comps
CV	349.7	349.4	339.9	341.6	339.2	339.6
adjCV	348.0	347.7	338.2	339.7	337.2	337.6

TRAINING: % variance explained

	1 comps	2 comps	3 comps	4 comps	5 comps	6 comps	7 comps	8 comps
X	38.31	60.16	70.84	79.03	84.29	88.63	92.26	94.96
Salary	40.63	41.58	42.17	43.22	44.90	46.48	46.69	46.75

	9 comps	10 comps	11 comps	12 comps	13 comps	14 comps	15 comps
X	96.28	97.26	97.98	98.65	99.15	99.47	99.75
Salary	46.86	47.76	47.82	47.85	48.10	50.40	50.55

	16 comps	17 comps	18 comps	19 comps
X	99.89	99.97	99.99	100.00
Salary	53.01	53.85	54.61	54.61

A few notes about the `summary()` of the `pcr()` results object. First, notice that there are two output panels. The first panel is the **Validation** and the second one is about **explained variance**.

Validation Output Panel

1. This panel shows CV **Square Root (MSE)** values, not MSE. Naturally, you can square these values to get the MSEs. The adjCV is a “bias-corrected” CV. It makes very little difference for our purposes, but adjCV makes some statistical adjustment that may come from sampling bias in CV testing. Let's pay more attention to the plain CV values, which are the real ones.
2. The CV RMSE goes down from 0 (Null model) to 1 component and it then goes up and down slightly until it gets to 6 components with a CV RMSE = 343.7. It then goes up and down again, with another low point at 16 components with a CV RMSE = 343.4. The absolute lowest point is at 18 components with a CV RMSE = 341.2.

Variance Explained Panel

1. The **X** row shows how much of the variance of the original predictors (i.e., X's) is explained by the PCs. For example, **6 PCs** explain **88.63%** of the variance in the predictors. The **1st PC** explains **38.31%** of the variance of the predictors. The model with **2 PCs** explains **60.16%** of the variance of the predictors. This means that the **2nd PC** alone explains $60.16 - 38.31 = 21.85\%$ of the variance of the predictors. As expected, the 1st PC explains more variance than the 2nd PC, and so on, but cumulatively, more PCs explain increasingly more variance.
2. **Salary** is the outcome variable, so its row shows how much of the outcome variable variance is explained by the model. For example, a model with **6 PCs** explain **46.48%** of the variance in Salary. This is equivalent to the 6 PC model's **R-Squared**.

**** Selecting the Best Model****

Which model is the best to select depends on the analytics goals:

1. If the analytics goal is predictive accuracy, it is best to select the model with the lowest CV RMSE or 18 PCs in this case.
2. If the analytics goal is interpretability, it is best to select the model with the largest possible number of PCs that has an acceptable CV RMSE, which in this case is also 18 PCs.
3. If the analytics goal is to substantially reduce the dimension of the model (this is often the goal when there is a very large number of predictors and not a lot of observations, which is typical in survey data analysis, for example), it is best to select the model with the smallest number of PC's that explains at least 70% of the variance of the predictors, which is 3 PCs in this case, which explains 70.84%.

PCR Loadings and Scores

The `pcr()` object is very complex and is full of information. You can explore this object's results with `str(pcr.fit)`. If you run this function, you will see an object attribute named **\$loadings**, which is a data frame containing one column for each component model and one row for each predictors. I'm only listing the first 8 PC models for simplicity, but you can remove the `[,1:8]` index to see all of them.

```
round(pcr.fit$loadings[,1:8],
      digits = 3)
```

	Comp 1	Comp 2	Comp 3	Comp 4	Comp 5	Comp 6	Comp 7	Comp 8
AtBat	0.198	0.384	-0.089	0.032	-0.028	-0.071	-0.107	0.270
Hits	0.196	0.377	-0.074	0.018	0.005	-0.082	-0.130	0.389
HmRun	0.204	0.237	0.216	-0.236	-0.078	-0.150	0.506	-0.226
Runs	0.198	0.378	0.017	-0.050	0.039	-0.137	-0.202	0.115
RBI	0.235	0.315	0.073	-0.139	-0.024	-0.112	0.319	0.005
Walks	0.209	0.230	-0.046	-0.131	0.032	-0.019	-0.558	-0.623
Years	0.283	-0.262	-0.035	0.095	0.010	0.033	0.012	0.138
CAtBat	0.330	-0.193	-0.084	0.091	-0.012	0.024	-0.012	0.147
CHits	0.331	-0.183	-0.086	0.084	-0.009	0.029	-0.020	0.195
CHmRun	0.319	-0.126	0.086	-0.074	-0.033	-0.041	0.229	-0.249
CRuns	0.338	-0.172	-0.053	0.069	0.018	0.007	-0.063	0.085
CRBI	0.340	-0.168	-0.015	0.007	-0.028	0.011	0.119	-0.002
CWalks	0.317	-0.192	-0.042	0.030	0.034	0.034	-0.178	-0.263
LeagueN	-0.054	-0.095	-0.548	-0.396	-0.012	-0.137	0.077	-0.026
DivisionW	-0.026	-0.037	0.016	0.043	-0.986	-0.091	-0.113	0.003
PutOuts	0.078	0.156	-0.051	-0.288	-0.106	0.924	0.019	0.065
Assists	-0.001	0.169	-0.398	0.524	0.011	0.035	0.013	-0.076
Errors	-0.008	0.201	-0.383	0.422	-0.055	0.148	0.373	-0.301
NewLeagueN	-0.042	-0.078	-0.545	-0.418	-0.014	-0.157	0.023	0.068

You can verify that the **sum of squared** component loadings for each row or each column always equals 1. These weights can be used to convert predictors into PCs and the PCs back into predictors:

```
sum(pcr.fit$loadings[1, ] ^ 2) # Sum of squared loadings for 1st predictor
```

```
[1] 1
```

```
sum(pcr.fit$loadings[, 1] ^ 2) # Sum of squared loadings for 1st PC
```

```
[1] 1
```

If you take the predicted values for some observations and convert them into their PC equivalent values, these values are called **PC Scores**. I'm only listing the first 6 scores of the first 10 observations, for simplicity:

```
round(pcr.fit$scores[1:10, 1:6],
      digits = 4)
```

	Comp 1	Comp 2	Comp 3	Comp 4	Comp 5	Comp 6
-Alan Ashby	-0.0096	-1.8670	-1.2627	-0.9337	-1.1075	1.2097
-Alvin Davis	0.4107	2.4248	0.9075	-0.2637	-1.2297	1.8231
-Andre Dawson	3.4602	-0.8244	-0.5544	-1.6136	0.8559	-1.0268
-Andres Galarraga	-2.5534	0.2305	-0.5187	-2.1721	0.8187	1.4889
-Alfredo Griffin	1.0257	1.5705	-1.3288	3.4874	-0.9816	0.5127
-Al Newman	-3.9731	-1.5044	0.1552	0.3691	1.2070	0.0334
-Argenis Salazar	-3.4452	-0.5988	0.6253	1.9960	-0.8055	0.2056
-Andres Thomas	-3.4258	-0.1133	-1.9959	0.7664	-1.0142	-0.2723
-Andre Thornton	3.8923	-1.9442	1.8170	-0.0267	1.1350	-0.8195
-Alan Trammell	3.1688	2.3878	-0.7930	2.5641	0.9455	-0.0611

PCR Coefficients

Reconstructing coefficients for the actual predictors from the PCs is straightforward. The **\$coefficients** attribute of the **pcr()** contains these values in a list.

To get the coefficients for the 1-PC model

```
pcr.fit$coefficients[, , 1]
```

	AtBat	Hits	HmRun	Runs	RBI	Walks
21.13207878	20.87321071	21.77988064	21.13705999	25.06279956	22.26529508	
	Years	CAtBat	CHits	CHmRun	CRuns	CRBI
30.11445915	35.21789413	35.24760132	33.99408860	36.04328244	36.27081015	
	CWalks	LeagueN	DivisionW	PutOuts	Assists	Errors
33.76212997	-5.80503669	-2.74157997	8.28029613	-0.08969488	-0.83758395	
	NewLeagueN					
-4.46643991						

```
# Or, alternatively: coef(pcr.fit, ncomp = 1)
```

To get the coefficients for the 3-PC model:

```
pcr.fit$coefficients[ , , 3]
```

AtBat	Hits	HmRun	Runs	RBI	Walks	Years
31.596172	30.841116	21.650526	28.894882	30.091792	28.345853	25.276570
CAtBat	CHits	CHmRun	CRuns	CRBI	CWalks	LeagueN
33.076799	33.388213	29.160504	33.604363	32.997441	30.624769	5.466047
DivisionW	PutOuts	Assists	Errors	NewLeagueN		
-3.929845	12.899211	13.245133	12.826601	7.109718		

```
# Or alternatively coef(pcr.fit, ncomp = 3)
```

To get coefficients for the second variable (Hits) of the 3-PC model

```
pcr.fit$coefficients[2, , 3]
```

```
[1] 30.84112
```

To get coefficients for, say 2- to 4-PC component models:

```
pcr.fit$coefficients[ , , 2:4]
```

	2 comps	3 comps	4 comps
AtBat	29.438966	31.596172	30.411575
Hits	29.039128	30.841116	30.174755
HmRun	26.912608	21.650526	30.389560
Runs	29.312723	28.894882	30.745533
RBI	31.870731	30.091792	35.242082
Walks	27.235049	28.345853	33.185988
Years	24.434849	25.276570	21.744656
CAtBat	31.042550	33.076799	29.700428
CHits	31.288812	33.388213	30.284689
CHmRun	31.260422	29.160504	31.912973
CRuns	32.314419	33.604363	31.040899
CRBI	32.632509	32.997441	32.750143
CWalks	29.599532	30.624769	29.499314
LeagueN	-7.865898	5.466047	20.140856
DivisionW	-3.535498	-3.929845	-5.513888
PutOuts	11.651168	12.899211	23.556308
Assists	3.560723	13.245133	-6.174373
Errors	3.507803	12.826601	-2.808121
NewLeagueN	-6.145712	7.109718	22.588848

```
# Or alternatively coef(pcr.fit, ncomp = 2:4), different format
```

To get coefficients for a group of models, say 3, 6 and 18 PCs

```
pcr.fit$coefficients[ , , c(3, 6, 18)]
```

	3 comps	6 comps	18 comps
AtBat	31.596172	24.363042	-287.1638712
Hits	30.841116	25.321422	330.3182702

HmRun	21.650526	16.517824	35.8569392
Runs	28.894882	24.483536	-55.7545172
RBI	30.091792	26.859813	-25.4323629
Walks	28.345853	33.873700	133.8275233
Years	25.276570	24.422920	-15.0311528
CAtBat	33.076799	30.534064	-425.9232643
CHits	33.388213	31.617866	151.1036646
CHmRun	29.160504	27.460383	-0.3535161
CRuns	33.604363	32.497526	452.9583268
CRBI	32.997441	31.827587	239.5044763
CWalks	30.624769	33.605665	-206.9835246
LeagueN	5.466047	10.910631	31.7984349
DivisionW	-3.929845	-68.868255	-58.5994581
PutOuts	12.899211	74.954304	78.7188346
Assists	13.245133	-3.328012	54.5749754
Errors	12.826601	3.191508	-22.7108234
NewLeagueN	7.109718	11.959825	-13.0025534

```
# Or alternatively coef(pcr.fit, ncomp = c(3, 6, 18))
```

Can you spot the most biased model and the most biased coefficient in that model? Look for any coefficient that changes substantially in sign and magnitude. The most unbiased model of the 3 is the model with 18 PCs, because it is the closest to the OLS. Notice that some coefficients change in value to some extent, but the change is not much (e.g., HmRun, LeagueN), but others change wildly and even change signs (e.g., AtBat, CWalks). Listing the coefficients for various PC models is useful because you can easily spot the most problematic predictors when it comes to bias and interpretation.

Predictions with PCR

To do predictions, you can use any data frame with values to feed to the `predict()` function. The values need to be in the same format as the original predictors. For this illustration I use 5% of the existing data to test our predictions. For this illustration, we use a model with 18 components for the prediction. Let's use a random subset of 5% of the data to make predictions using the trained model with the test data:

```
set.seed(2) # Randomly selected
test <- sample(1:nrow(Hitters),
              0.05 * nrow(Hitters))
Hitters.test <- Hitters[test, ] # Test sub-sample
pcr.pred <- predict(pcr.fit,
                  Hitters.test,
                  ncomp = 18)

pcr.pred
```

, , 18 comps

	Salary
-Rick Dempsey	356.2034
-Willie Upshaw	1022.4237
-Rick Leach	247.8984
-Mookie Wilson	649.1189
-Eric Davis	542.5348
-John Shelby	197.9010
-Dale Murphy	1041.7816
-Kevin Bass	564.6475


```
-Tony Bernazard    797.5811
-Bill Schroeder    208.4064
-Doug DeCinces     694.5173
-Herm Winningham   211.7767
-Jose Cruz         872.0351
```

We can then compute the MSE for these predictions:

```
MSE <- round(
  mean( (pcr.pred - Hitters.test$Salary) ^ 2 ),
  digits = 2)
paste("MSE =", MSE)
```

```
[1] "MSE = 73830.35"
```

And also the RMSE:


```
RMSE <- round(sqrt(MSE), digits = 2)
paste("RMSE =", RMSE)
```

```
[1] "RMSE = 271.72"
```

Partial Least Squares (PLS) Regression

The process for fitting **PLSR** models is the same as for **PCR** models, except that we use `pls()` function from the same `{pls}` library instead of `pcr()`. The interpretation and analysis of the outputs is identical to **PCR**. The only thing that changes is how the PCs are constructed, but this is not visible in the **PLSR** output. In the next script lines, repeat the same steps above, but with the `pls()` function.

```
library(pls)
library(ISLR) # Has the Hitters data set
set.seed(2) # Arbitrary
pls.fit <- pls(Salary ~ .,
  data = Hitters,
  scale = T,
  validation = "CV")
validationplot(pls.fit,
  val.type = "MSEP",
  legendpos = "topright")
```



To get the RMSE instead of the MSE above, just change the validation type to `val.type = "RMSEP"`.

Now let's display the summary of the PLSR model output

```
summary(pls.fit)
```

Data: X dimension: 263 19

Y dimension: 263 1

Fit method: kernelpls

Number of components considered: 19

VALIDATION: RMSEP

Cross-validated using 10 random segments.

	(Intercept)	1 comps	2 comps	3 comps	4 comps	5 comps	6 comps
CV	452	349.2	348.4	348.2	348.1	353.0	350.6
adjCV	452	348.8	347.6	347.1	346.9	350.9	348.3
	7 comps	8 comps	9 comps	10 comps	11 comps	12 comps	13 comps
CV	345.0	340.9	342.8	339.5	339.3	338.3	340.1
adjCV	343.2	339.1	340.7	337.9	337.6	336.6	338.2
	14 comps	15 comps	16 comps	17 comps	18 comps	19 comps	
CV	340.5	340	339.6	339.2	339.3	339.6	
adjCV	338.5	338	337.6	337.3	337.4	337.6	

TRAINING: % variance explained

	1 comps	2 comps	3 comps	4 comps	5 comps	6 comps	7 comps	8 comps
X	38.08	51.03	65.98	73.93	78.63	84.26	88.17	90.12
Salary	43.05	46.40	47.72	48.71	50.53	51.66	52.34	53.26
	9 comps	10 comps	11 comps	12 comps	13 comps	14 comps	15 comps	
X	92.92	95.00	96.68	97.68	98.22	98.55	98.98	

Salary	53.52	53.77	54.04	54.20	54.32	54.47	54.54
	16 comps	17 comps	18 comps	19 comps			
X	99.24	99.71	99.99	100.00			
Salary	54.59	54.61	54.61	54.61			

Notice in the output the similarities and differences between PLSR and PCR. With PLSR, the first marked elbow is also at 1 PC. But the next interesting elbow is at 8 PCs, not 6. But then, the CV MSE goes down slightly and remains relatively flat. These are the recommended models from this output:

1. If the analysis goal is predictive accuracy, it would be best to select the model with the lowest CV RMSE or 17 PCs in this case, with CV RMSE = 339.2.
2. If the analysis goal is interpretability, it would be best to select the largest possible model with an acceptable CV RMSE, which in this case is also 17 PCs. However, since the difference in CV RMSE between 17 PCs and the full 19 PCs is minuscule, it would be OK to interpret the results with the 19 PC model.
3. If the analysis goal is to substantially reduce the dimension of the model, it would be best to select the model with the smallest number of PC's that explain at least 70% of the variance of the predictors, or 4 PCs in this case, which explains 73.93% of the variance in the predictors, which is a good representation of the data.

PLSR Loadings and Scores

PLSR PC Loadings

```
round(pls.fit$loadings[,1:8],
      digits = 3)
```

	Comp 1	Comp 2	Comp 3	Comp 4	Comp 5	Comp 6	Comp 7	Comp 8
AtBat	0.226	0.347	-0.396	0.099	0.109	-0.092	-0.199	-0.323
Hits	0.223	0.356	-0.371	0.147	0.178	-0.016	-0.108	-0.167
HmRun	0.218	0.091	-0.308	-0.014	-0.652	0.389	0.170	0.131
Runs	0.225	0.337	-0.375	0.164	-0.079	0.080	-0.036	-0.077
RBI	0.257	0.231	-0.343	0.043	-0.268	0.240	0.128	-0.008
Walks	0.229	0.283	-0.168	0.029	-0.098	0.043	-0.050	0.317
Years	0.266	-0.344	0.266	-0.166	-0.087	-0.240	-0.005	0.123
CAtBat	0.320	-0.258	0.213	-0.121	0.199	-0.144	0.012	0.036
CHits	0.321	-0.238	0.211	-0.103	0.243	-0.116	0.029	0.088
CHmRun	0.311	-0.226	0.135	-0.025	0.020	0.252	0.068	-0.257
CRuns	0.329	-0.234	0.203	-0.073	0.170	-0.089	0.034	0.015
CRBI	0.331	-0.241	0.192	-0.074	0.179	0.066	0.051	-0.058
CWalks	0.307	-0.233	0.230	-0.124	-0.042	-0.152	-0.041	-0.020
LeagueN	-0.046	0.297	0.209	-0.860	0.421	0.102	0.033	-0.069
DivisionW	-0.040	-0.376	-0.252	-0.293	0.491	0.164	-0.932	0.599
PutOuts	0.100	0.467	0.095	0.236	-0.031	-0.044	-0.737	0.746
Assists	0.010	0.131	-0.236	-0.012	0.916	-0.654	0.324	-0.244
Errors	0.005	0.165	-0.289	-0.156	0.565	-0.626	0.522	0.305
NewLeagueN	-0.032	0.313	0.188	-0.890	0.299	0.058	-0.068	-0.076

PLSR PC Scores

```
round(pls.fit$scores[1:10, 1:6], digits = 4)
```

	Comp 1	Comp 2	Comp 3	Comp 4	Comp 5	Comp 6
-Alan Ashby	-0.1090	-0.0879	1.1147	-1.4059	-0.6158	-1.2286
-Alvin Davis	0.6671	0.8786	-1.0206	0.9639	0.0307	0.1497
-Andre Dawson	3.4717	0.5270	1.2976	-0.3869	0.6279	2.0307
-Andres Galarraga	-2.1299	2.4542	2.0764	0.2078	-0.1079	0.5837
-Alfredo Griffin	0.9771	-0.7937	-2.1395	0.4122	0.8415	-2.3038
-Al Newman	-4.0037	0.1500	1.6439	0.7911	0.5309	0.5485
-Argenis Salazar	-3.6685	-1.3440	-0.5804	0.9552	0.7630	-0.4132
-Andres Thomas	-3.4262	-0.3027	-1.0077	-1.1485	0.7244	-0.4149
-Andre Thornton	3.5184	-1.3746	1.0514	0.3408	-0.8714	0.0946
-Alan Trammell	3.2932	0.1716	-1.7483	0.4770	0.1101	-1.6399

PLSR Coefficients

To get the coefficients for the 1 PC-model

```
pls.fit$coefficients[ , 1]
```

	AtBat	Hits	HmRun	Runs	RBI	Walks
25.0420570	27.8270677	21.7597795	26.6334747	28.5110396	28.1564522	
Years	CAAtBat	CHits	CHmRun	CRuns	CRBI	
25.4154350	33.3750764	34.8197471	33.2986538	35.6931216	35.9651267	
CWalks	LeagueN	DivisionW	PutOuts	Assists	Errors	
31.0715657	-0.9059591	-12.2120349	19.0607903	1.6135259	-0.3425902	
NewLeagueN						
-0.1798022						

```
# Or, alternatively: coef(pls.fit, ncomp = 1)
```

To get the coefficients for the 3 PC-model:

```
pls.fit$coefficients[ , 3]
```

	AtBat	Hits	HmRun	Runs	RBI	Walks
11.5612560	43.0738184	-3.0788552	29.4670885	21.0426670	43.6363824	
Years	CAAtBat	CHits	CHmRun	CRuns	CRBI	
5.4373897	28.9571188	41.9403523	35.6444168	42.6088887	43.5878120	
CWalks	LeagueN	DivisionW	PutOuts	Assists	Errors	
17.0901059	25.0699041	-69.4031815	74.5802997	0.8434621	-16.4113867	
NewLeagueN						
17.3645626						

```
# Or alternatively coef(pls.fit, ncomp = 3)
```

To get the coefficients for the second variable (Hits) of the 3-PC model

```
pls.fit$coefficients[2, , 3]
```

```
[1] 43.07382
```

To get the coefficients for, say 2 to 4 PCR component models:

```
pls.fit$coefficients[ , , 2:4]
```

	2 comps	3 comps	4 comps
AtBat	26.988567	11.5612560	-6.652720
Hits	43.689562	43.0738184	58.942213
HmRun	12.561740	-3.0788552	-18.405673
Runs	36.166065	29.4670885	31.456684
RBI	30.845058	21.0426670	16.922246
Walks	42.533221	43.6363824	47.896733
Years	8.957764	5.4373897	-14.209470
CAtBat	26.430999	28.9571188	25.046713
CHits	33.829611	41.9403523	51.191402
CHmRun	30.279265	35.6444168	44.690717
CRuns	34.715277	42.6088887	51.881789
CRBI	35.195944	43.5878120	55.348469
CWalks	20.084235	17.0901059	-3.215364
LeagueN	17.985777	25.0699041	9.757598
DivisionW	-48.032748	-69.4031815	-81.335717
PutOuts	56.282325	74.5802997	89.176182
Assists	4.199492	0.8434621	8.658344
Errors	-4.327875	-16.4113867	-26.720518
NewLeagueN	15.096458	17.3645626	-8.808971

```
# Or alternatively coef(pls.fit, ncomp = 2:4) # different format
```

To get the coefficients for a group of models, say 1, 8 and 17 PCs

```
pls.fit$coefficients[ , , c(1, 8, 17)]
```

	1 comps	8 comps	17 comps
AtBat	25.0420570	-268.57609	-295.439407
Hits	27.8270677	208.90183	333.458425
HmRun	21.7597795	-10.42637	31.880957
Runs	26.6334747	28.55779	-51.938191
RBI	28.5110396	26.94159	-19.261427
Walks	28.1564522	131.95382	131.287369
Years	25.4154350	-99.56742	-16.399659
CAtBat	33.3750764	-23.34822	-409.837439
CHits	34.8197471	157.73966	124.930500
CHmRun	33.2986538	61.17685	-1.550428
CRuns	35.6931216	160.81868	464.859131
CRBI	35.9651267	136.84026	238.736301
CWalks	31.0715657	-197.27214	-205.460321
LeagueN	-0.9059591	42.70521	30.749788
DivisionW	-12.2120349	-61.70447	-57.865552
PutOuts	19.0607903	83.68301	79.361453
Assists	1.6135259	39.12289	54.817784
Errors	-0.3425902	-13.76418	-22.438269
NewLeagueN	-0.1798022	-31.85963	-11.533475

```
# Or alternatively coef(pls.fit, ncomp = c(1, 8, 17))
```

As with PCR, you can spot the most biased model and the most biased coefficient in that model. Interestingly, the models don't get as biased as with PCR when we reduce the number of components. Notice that some coefficients in the 8 PC model are different than the 17 PC model, but except for a few cases, not by much.

Predictions with PLSR

Again, the process is the same as for PCR:

```
set.seed(2) # Arbitrary
test <- sample(1:nrow(Hitters),
              0.05 * nrow(Hitters))
Hitters.test <- Hitters[test, ] # Test sub-sample
pls.pred <- predict(pls.fit,
                  Hitters.test,
                  ncomp = 17)

pls.pred
```

, , 17 comps

	Salary
-Rick Dempsey	349.9180
-Willie Upshaw	1022.1080
-Rick Leach	251.5487
-Mookie Wilson	652.8297
-Eric Davis	545.9409
-John Shelby	198.3570
-Dale Murphy	1038.4198
-Kevin Bass	566.6128
-Tony Bernazard	801.7325
-Bill Schroeder	207.5167
-Doug DeCinces	695.9014
-Herm Winningham	209.7285
-Jose Cruz	869.5758

We can then compute the MSE for these predictions as follows:

```
MSE <- round(
  mean( (pls.pred - Hitters.test$Salary) ^ 2 ),
  digits = 2)
paste("MSE =", MSE)
```

[1] "MSE = 74394.11"

And the RMSE by taking the square root:

```
RMSE <- round(sqrt(MSE), digits = 2)
paste("RMSE =", RMSE)
```

[1] "RMSE = 272.75"

Logistic and GLM Regressions with PCR and PLSR

Running PCR and PLSR models with classification models like Logistic or even other GLM models, is similar to the quantitative methods we discussed above. The only problem is that the `{pls}` R package does not work with **GLM** models. But there are many libraries to estimate these models. In the example below I illustrate estimating a **Logistic PCR** model using the `pcr()` function from the `{compositional}` R package.

As with the Ridge regression I illustrate how to setup a Logistic model with PCR or PLSR. Again, we can access the `SAHeart.data` dataset in Prof. Robert Tibshirani's web site, which contains variables associated with coronary heart disease. The dataset and variable descriptions are available at: <https://rdrr.io/cran/ElemStatLearn/man/SAheart.html>.

Once you have downloaded the data set, you can open it directly. I downloaded the file in CSV format with the name `Heart.csv`, which can be read as follows:

```
heart <- read.table("Heart.csv", sep = ",", head = T)
```

Once you have the data in the heart data frame object, the next step is to create the predictor matrix `x` and the outcome vector `y`. The outcome variable is binary, `chd = 1` for disease, `0` otherwise)

```
x <- model.matrix(chd ~ ., data = heart)[-1]
y <- heart$chd
```


We now can estimate a **Logistic PCR** model using the `pcr()` function from the `{compositional}` R package.

```
library(Compositional)
```

Technical Note: the outcome variable in the data set must be an integer. If the outcome is 0 or 1 as an integer, the function will fit a Logistic model. If it contains multiple integer values, it will fit a count data model.

The next step is to **tune** the model using the `pcr.tune(y, x, graph=T)` function. The default CV tuning method is **10FCV**, but you can use the `nfolds =` parameter to change the number of CV folds. The function renders a scree plot and also the mean deviance for every fold in each PC

```
heart.tune <- pcr.tune(y, x,
                      graph = T)
```



```
round(heart.tune$msp, digits = 4)
```

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]
[1,]	0.2318	0.2453	0.2427	0.2264	0.2077	0.2067	0.2100	0.2066	0.2067
[2,]	0.1643	0.1691	0.1584	0.1397	0.1546	0.1571	0.1540	0.1527	0.1539
[3,]	0.1958	0.1967	0.2110	0.2037	0.1989	0.1987	0.2127	0.2175	0.2183
[4,]	0.2135	0.2131	0.2201	0.2178	0.1894	0.1834	0.1911	0.1896	0.1898
[5,]	0.1372	0.1344	0.1341	0.1312	0.1242	0.1252	0.1242	0.1226	0.1227
[6,]	0.2263	0.2268	0.2170	0.2188	0.2126	0.2141	0.2173	0.2152	0.2182
[7,]	0.1577	0.1531	0.1568	0.1447	0.1351	0.1347	0.1314	0.1315	0.1315
[8,]	0.2479	0.2442	0.2360	0.2194	0.2132	0.2158	0.2151	0.2245	0.2251
[9,]	0.2046	0.2056	0.1989	0.1905	0.1787	0.1765	0.1763	0.1738	0.1755
[10,]	0.1853	0.1857	0.1743	0.1858	0.1817	0.1832	0.1814	0.1828	0.1832

To find the number of components that minimizes the CV deviance, we use the **\$k** attribute:

```
best.comp <- heart.tune$k
best.comp # Check it out
```

PC 6
6

You can now fit a PCR model with k components

```
heart.pcr <- pcr(y, x,
                 k = best.comp)
heart.pcr$be # Log-Odds coefficients
```


	PC6
sbp	0.02784825
tobacco	0.09741489
ldl	0.04965203
adiposity	0.01653809
famhistPresent	0.09514942
typea	0.05633556
obesity	-0.02619153
alcohol	-0.01738710
age	0.08177750

Let's display the results with log-odds and odds

```
results <- round(cbind(heart.pcr$be,
                      exp(heart.pcr$be)),
                digits = 4)
colnames(results) <- c("5 PC Log-Odds",
                      "5 PC Odds")
results
```

	5 PC Log-Odds	5 PC Odds
sbp	0.0278	1.0282
tobacco	0.0974	1.1023
ldl	0.0497	1.0509
adiposity	0.0165	1.0167
famhistPresent	0.0951	1.0998
typea	0.0563	1.0580
obesity	-0.0262	0.9741
alcohol	-0.0174	0.9828
age	0.0818	1.0852